

Efficient VM with JIT in Go

quasilyte @ GoWayFest 4.0 (2020)



Part 0/7

Backstory

- > 0 - Backstory
- 1 - go-jdk overview
- 2 - Making the code run fast
- 3 - GC-friendly slots
- 4 - Interop / FFI
- 5 - Object layout / mem alloc
- 6 - Challenges & limitations
- 7 - Closing words



Once upon a time:

“Can we use Lucene from Go?”

Once upon a time:

“Can we use Lucene from Go?”

Sure...

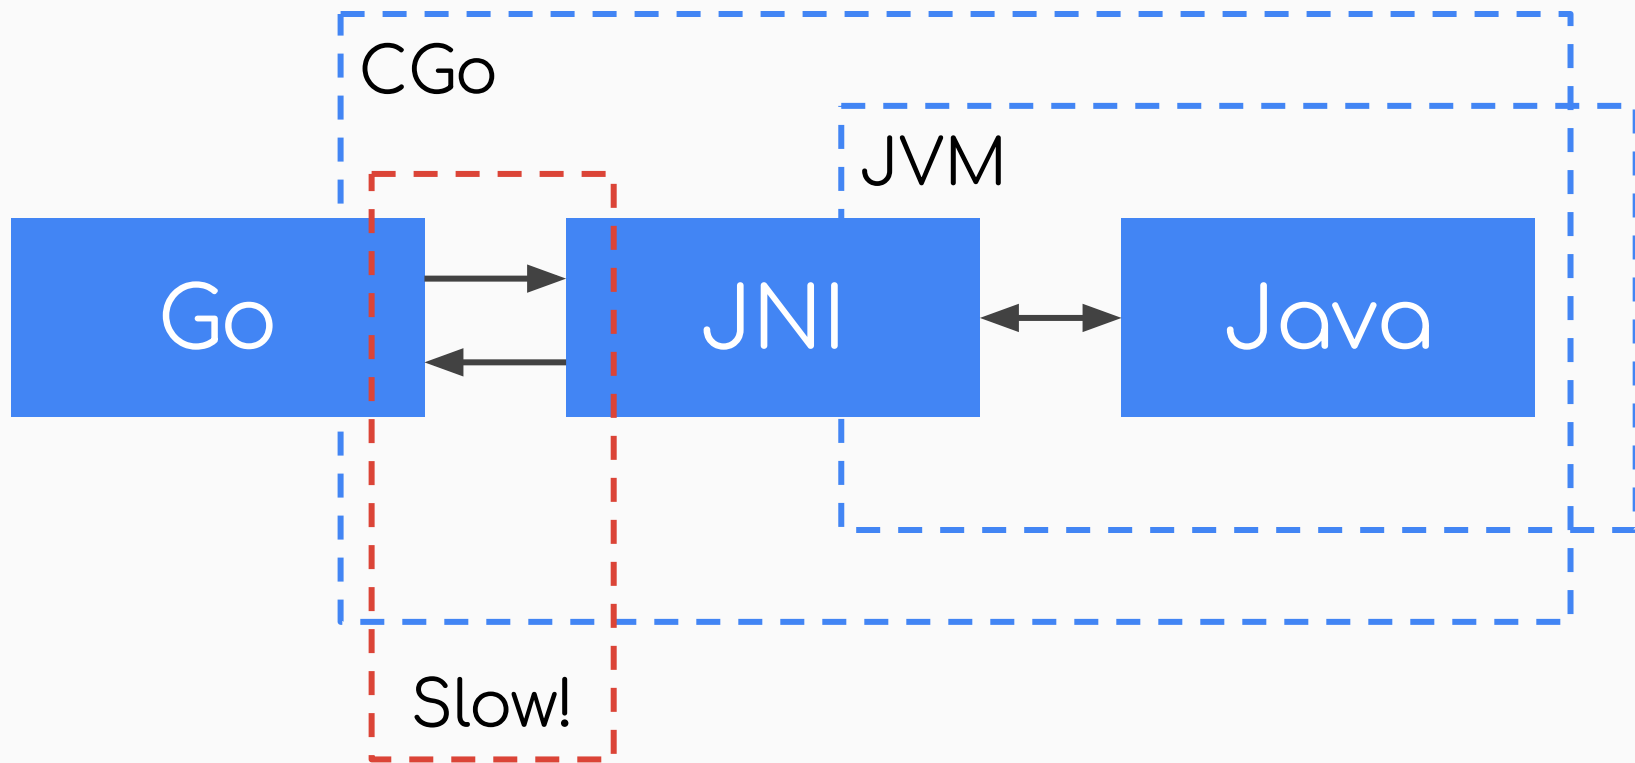
Once upon a time:
“Can we use Lucene from Go?”
Sure...



How to use Java from Go?

- JNI (with CGo blessing)
- Pass arguments through serialization

<https://github.com/timob/jnigi>



Why JNI is not good for Go?


- Locked OS thread for JVM goroutines

Why JNI is not good for Go?


- Locked OS thread for JVM goroutines
- Every JNI call has CGo call overhead

Why JNI is not good for Go?


- Locked OS thread for JVM goroutines
- Every JNI call has CGo call overhead
- Expensive Go↔JNI values conversion



Two active runtimes in one application



Two active runtimes in one application



Two active runtimes in one application

Long story short...

We're now using Lucene from our
Go application, but

Long story short...

We're now using Lucene from our
Go application, but
it bothers me how inefficient it is.
Can we do better?

Part 1/7

go-jdk overview

- 0 - Backstory
- > 1 - go-jdk overview
- 2 - Making the code run fast
- 3 - GC-friendly slots
- 4 - Interop / FFI
- 5 - Object layout / mem alloc
- 6 - Challenges & limitations
- 7 - Closing words



Let's try build an efficient JVM
that can be easily embedded into
Go applications.

Me (just now)

Quote

So, what exactly do we want?

- Cheap Go↔Java calls (and no CGo)

So, what exactly do we want?

- Cheap Go↔Java calls (and no CGo)
- Optimized machine code (no interpretation)

So, what exactly do we want?


- Cheap Go↔Java calls (and no CGo)
- Optimized machine code (no interpretation)
- Efficient objects layout and allocation

So, what do we want?

- Cheap (e.g. C, C++)
- Optimized (e.g. C, C++)
- Efficient (e.g. C, C++)

DO IT

Direct connection



go-jdk project

- Java class file loader
- JIT compiler (non-tracing)
- Runtime and interop primitives
- Utility tools like “javap”

<https://github.com/quasilyte/go-jdk>



Source code



Class file



JVM

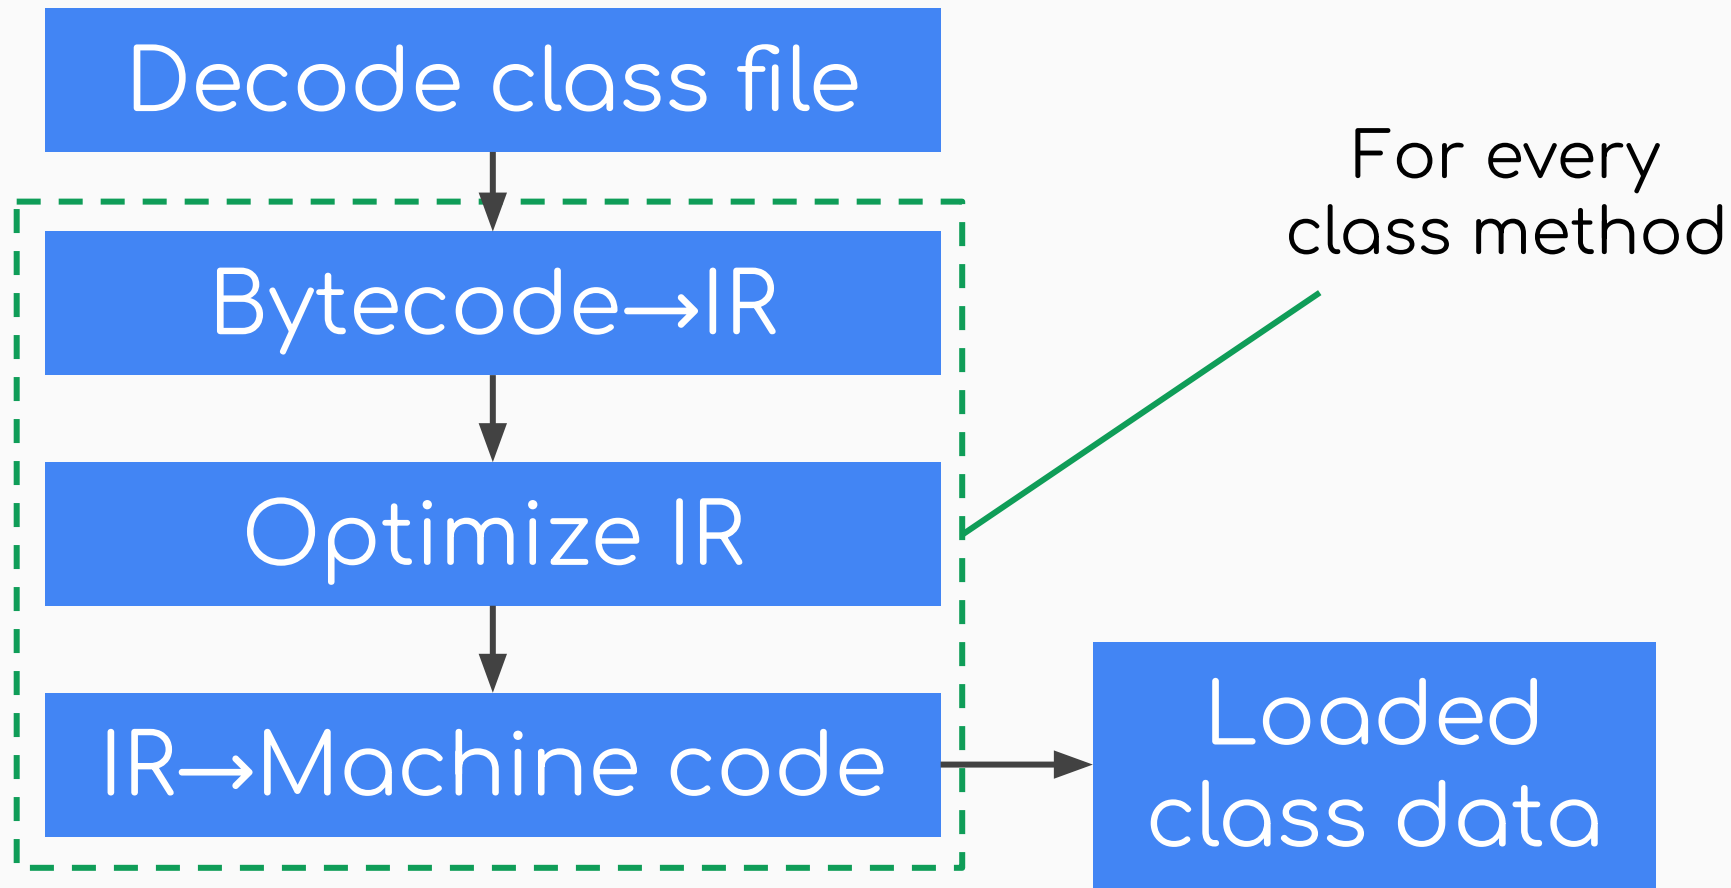
Source code

Class file

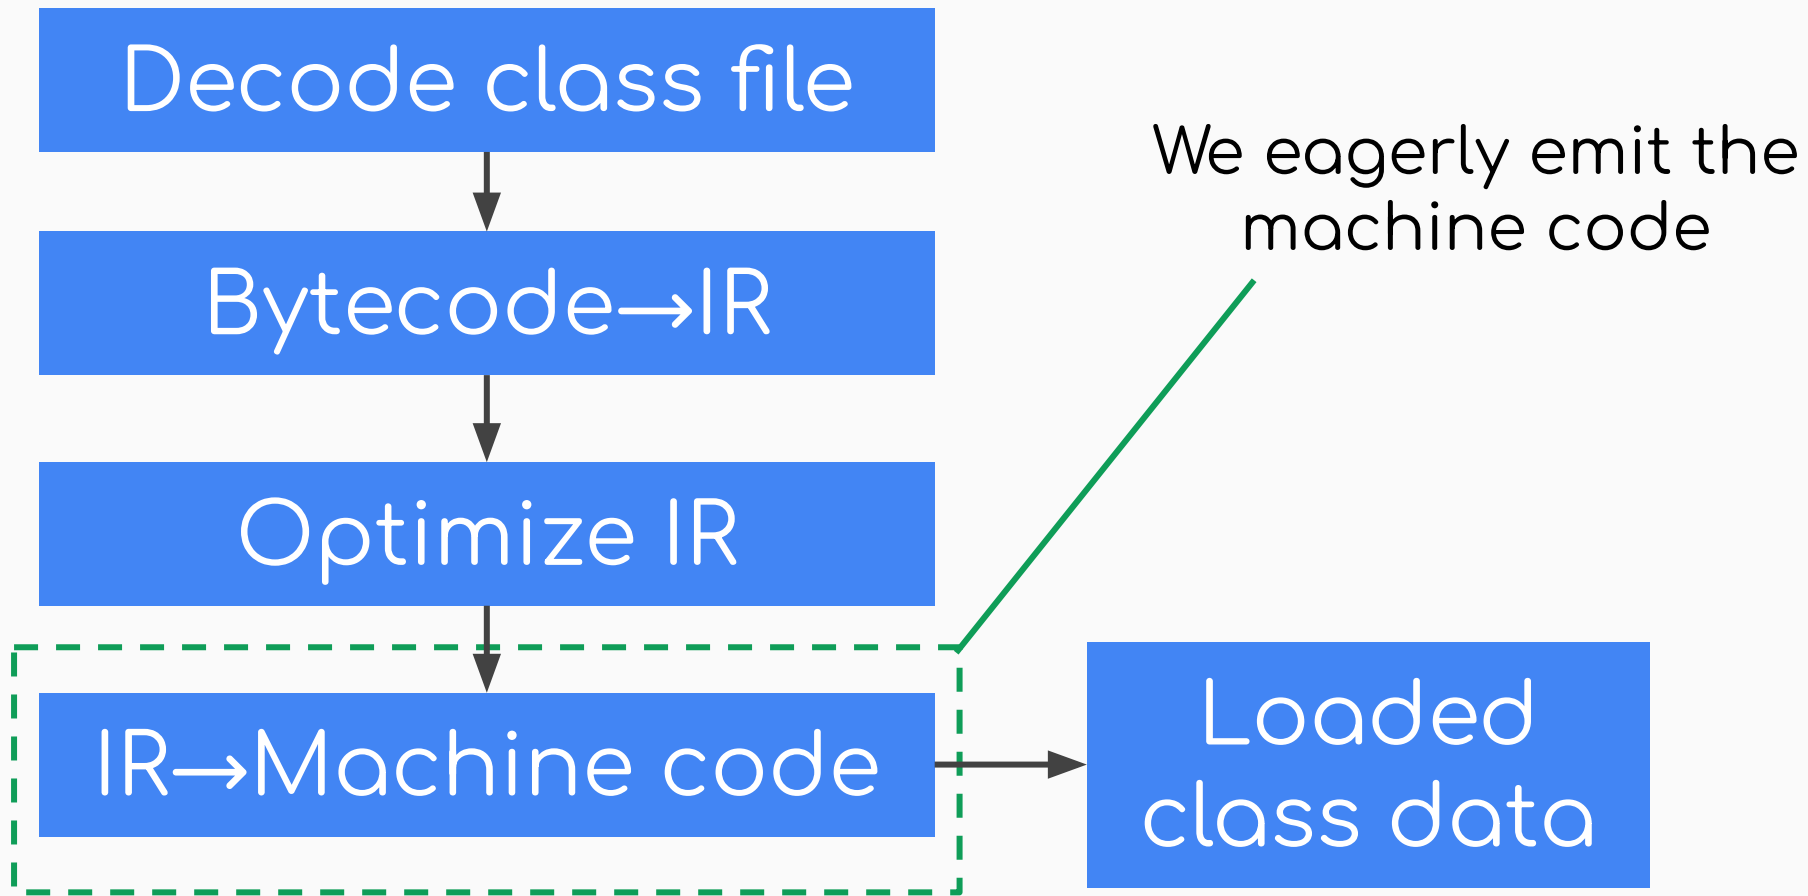
JVM

go-jdk uses class
files as its input

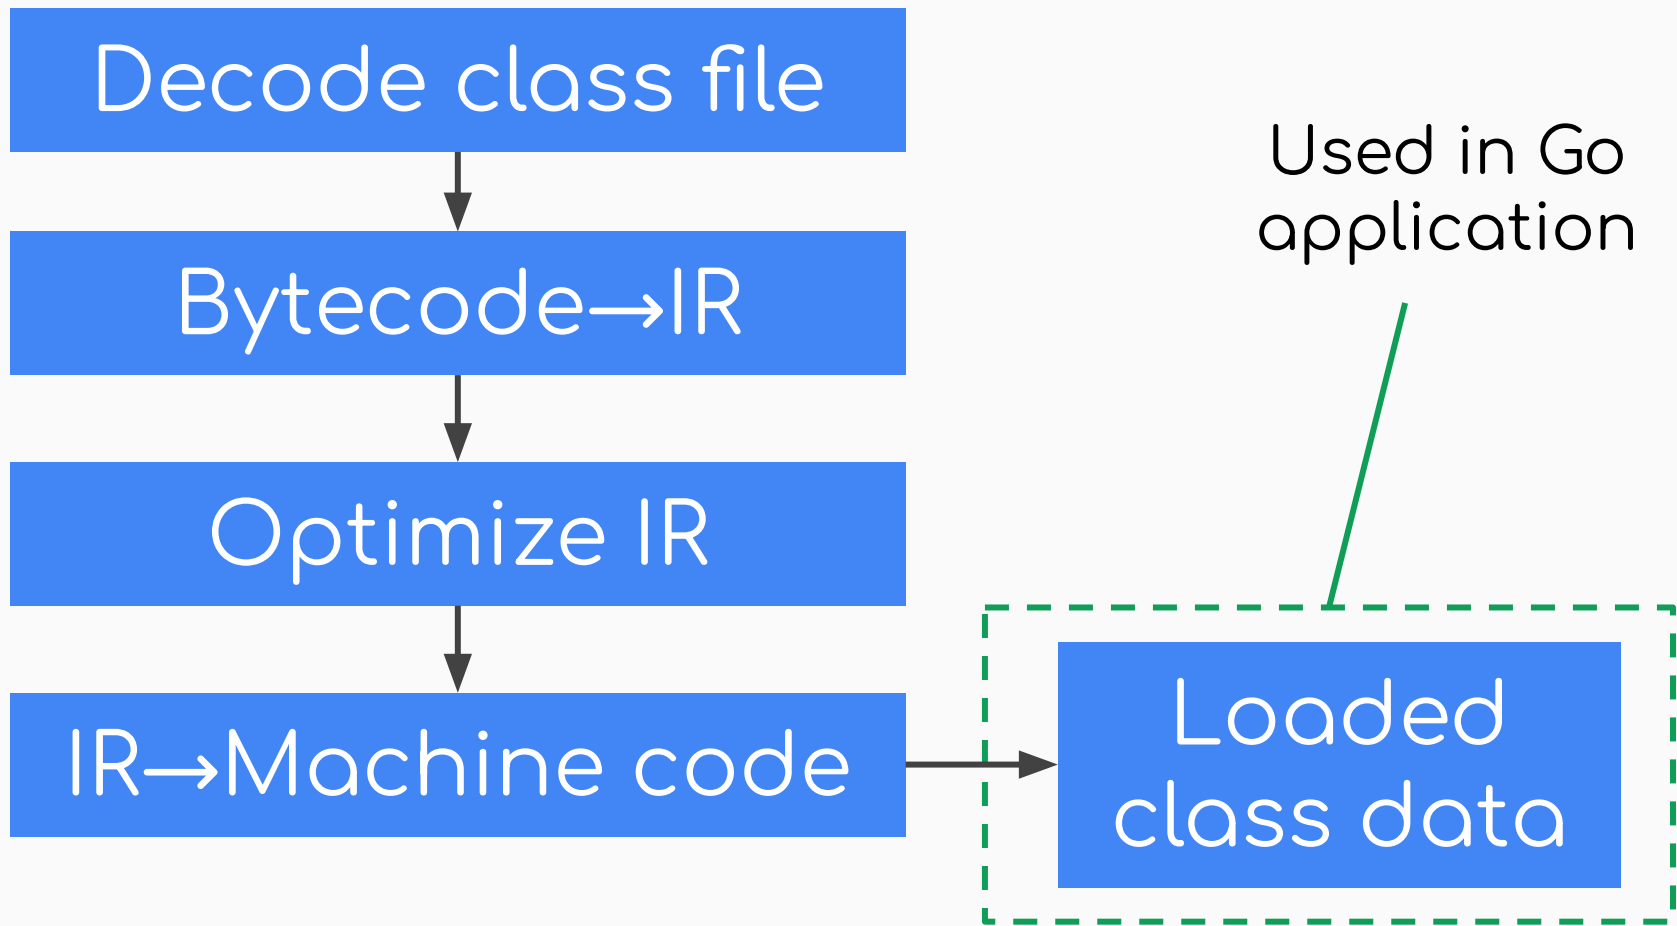
go-jdk inputs



How class is loaded



How class is loaded





How class is loaded

Part 2/7

Making the code run fast

- 0 - Backstory
- 1 - go-jdk overview
- > 2 - Making the code run fast
- 3 - GC-friendly slots
- 4 - Interop / FFI
- 5 - Object layout / mem alloc
- 6 - Challenges & limitations
- 7 - Closing words





Convert Java class file into... what?



Idea 1: direct bytecode to machine code translation

Our example class and example static method

```
class Example {  
    public static int add1(int x) {  
        return x + 1;  
    }  
}
```


bytecode→amd64

iload_0

MOVQ local_0(CX), AX

MOVQ AX, (CX)

ADDQ \$8, CX

iconst_1

MOVQ \$1, (CX)

ADDQ \$8, CX

iadd

MOVQ -16(CX), AX

ADDQ -8(CX), AX

MOVQ AX, -16(CX)

SUBQ \$8, CX

ireturn

RET

bytecode→amd64

iload_0

MOVQ local_0(CX), AX

MOVQ AX, (CX)

ADDQ \$8, CX

iconst_1

MOVQ \$1, (CX)

ADDQ \$8, CX

iadd

MOVQ -16(CX), AX

ADDQ -8(CX), AX

MOVQ AX, -16(CX)

SUBQ \$8, CX

ireturn

RET

Stack
bookkeeping



bytecode→amd64

iload_0

iconst_1

iadd

ireturn

MOVQ local_0(CX), AX

MOVQ AX, (CX)

ADDQ \$8, CX

MOVQ \$1, (CX)

ADDQ \$8, CX

MOVQ -16(CX), AX

ADDQ -8(CX), AX

MOVQ AX, -16(CX)

SUBQ \$8, CX

RET


Hard to
analyze and
optimize

Stack vs Register architecture

Suggested reading:

[VM Showdown: Stack Versus Registers](#)

We can't change the input bytecode format,
but we can add intermediate representation.



Idea 2: add intermediate representation

bytecode→IR

```
iload_0  
iconst_1  
iadd
```

```
ireturn
```

```
r1 = iadd r0 1
```

```
iret r1
```

IR→amd64

```
r1 = iadd r0 1
```

```
iret r1
```

```
MOVQ local_0(CX), AX  
ADDQ $1, AX  
MOVQ AX, local_1(CX)
```

```
MOVQ local_1, AX  
RET
```

IR→amd64

```
r1 = iadd r0 1
```

```
iret r1
```

```
MOVQ local_0(CX), AX
```

```
ADDQ $1, AX
```

```
MOVQ AX, local_1(CX)
```

```
MOVQ local_1, AX
```

```
RET
```

Can be optimized-out

IR→amd64

```
[ret] = iadd r0 1
```

```
iret [ret]
```

Mapped to
AX (or X0)

```
MOVQ local_0(CX), AX  
ADDQ $1, AX
```

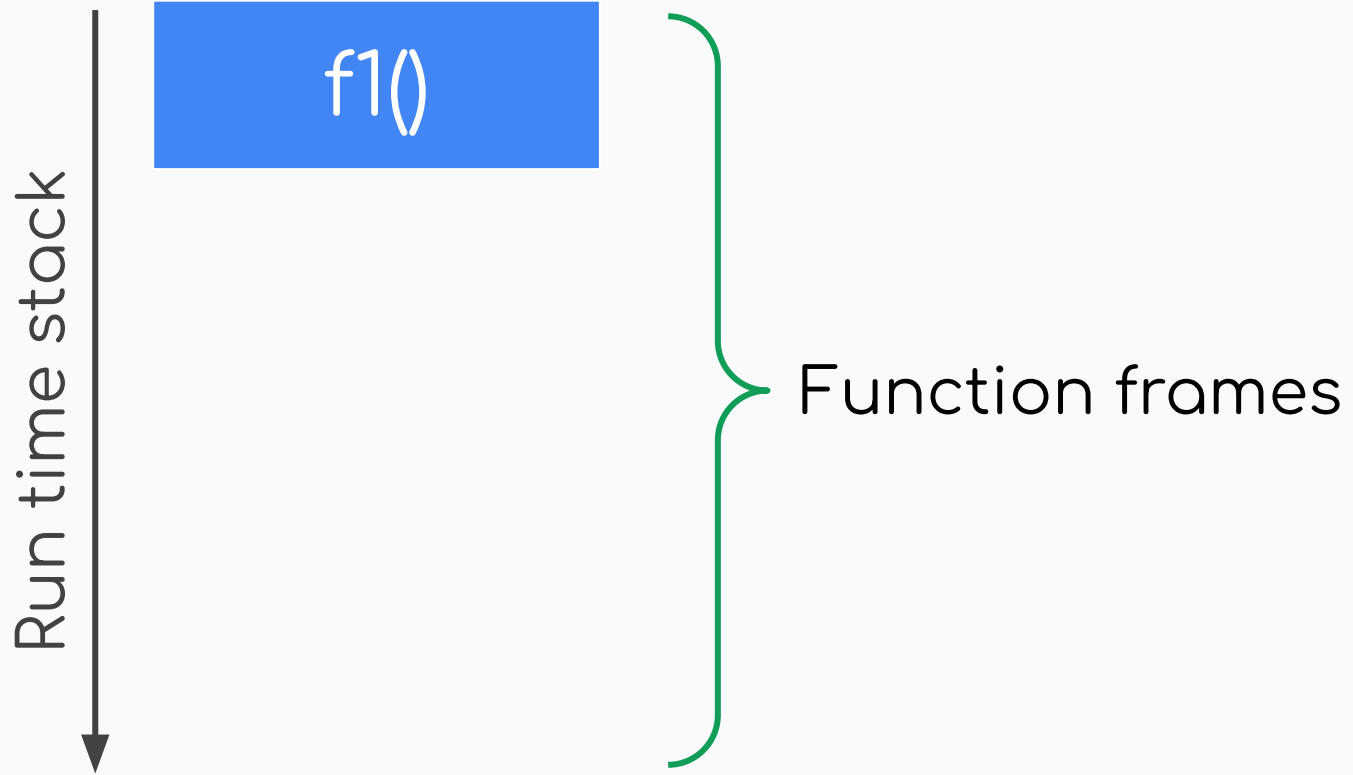
```
RET
```

Part 3/7

GC-friendly slots

- 0 - Backstory
- 1 - go-jdk overview
- 2 - Making the code run fast
- > 3 - GC-friendly slots
- 4 - Interop / FFI
- 5 - Object layout / mem alloc
- 6 - Challenges & limitations
- 7 - Closing words





Run time data lives inside the run time stack



`f1() calls f2()`

Run time data lives inside the run time stack



f1() calls f2()
f2() calls f3()

Run time data lives inside the run time stack



f1() calls f2()
f2() calls f3()
f3() returns


Run time data lives inside the run time stack

Run time stack
↓

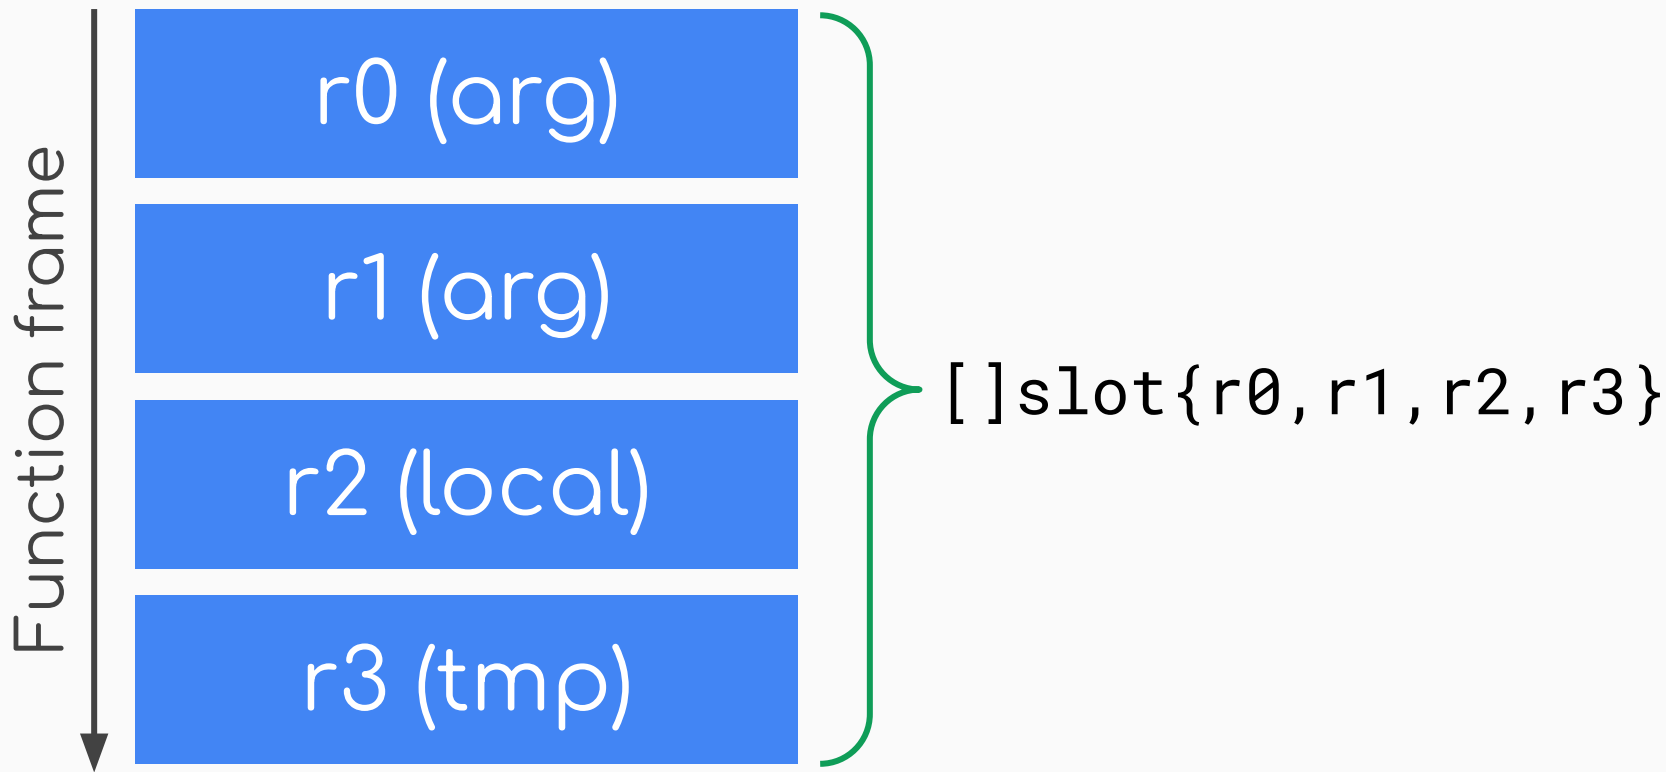
f1()

```
f1() calls f2()  
f2() calls f3()  
f3() returns  
f2() returns
```


Run time data lives inside the run time stack




Function frame model (abstract)



Function frame model (concrete)




What do we store inside a slot?



What do we store inside a slot?

Seems like
everything fits in
64-bit slots



What do we store inside a slot?

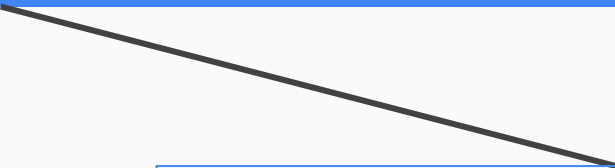
Function frame

r0

r1

r2

...



```
type slot struct {  
    value uint64  
}
```

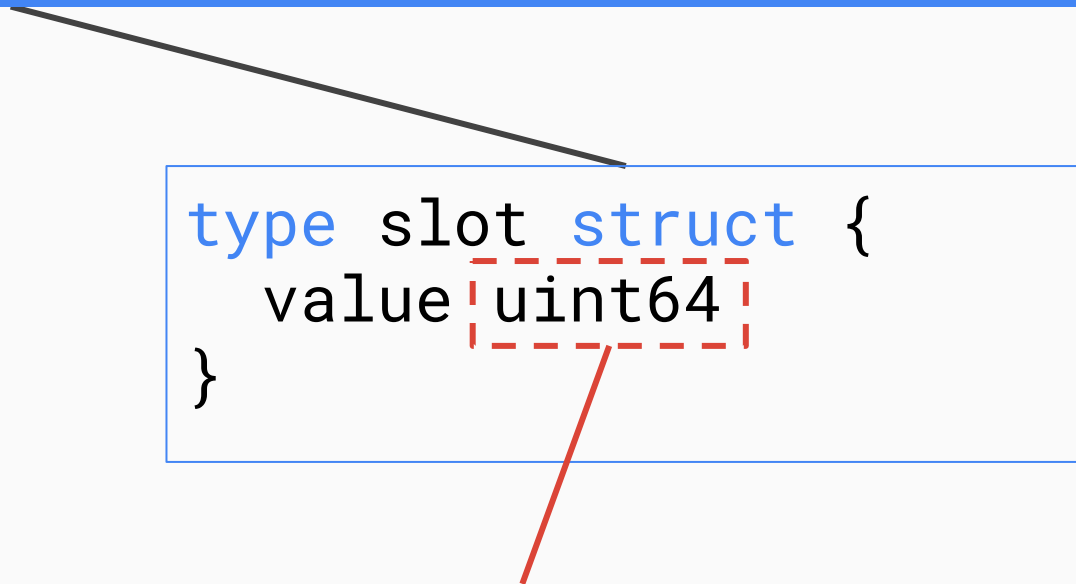
Uint64 slots

r0

r1

r2

...



```
type slot struct {  
    value uint64  
}
```

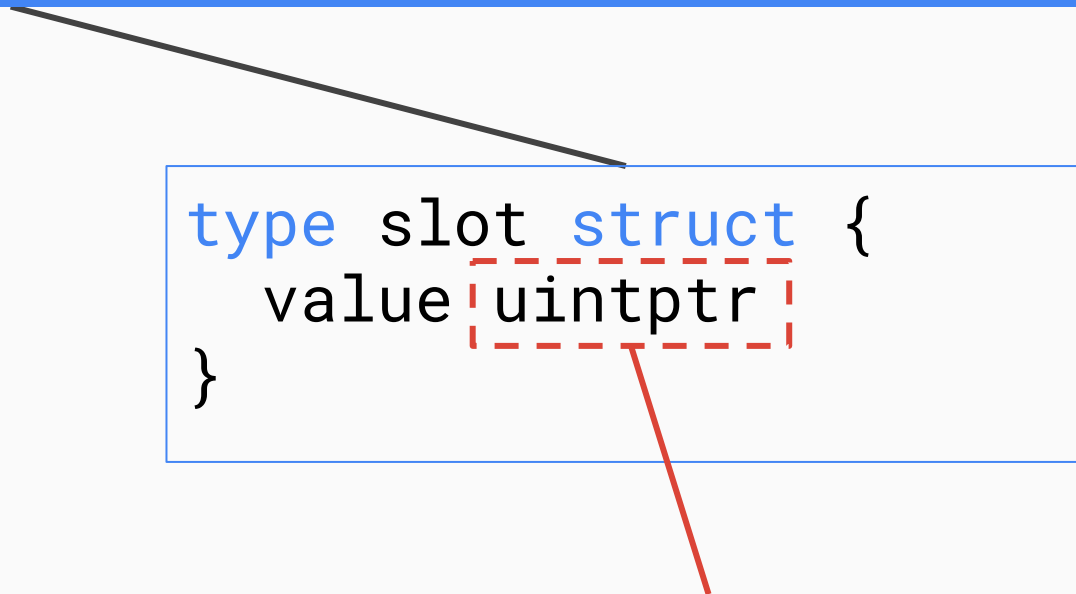
Not safe to store **pointers** there!

r0

r1

r2

...



```
type slot struct {  
    value uintptr  
}
```

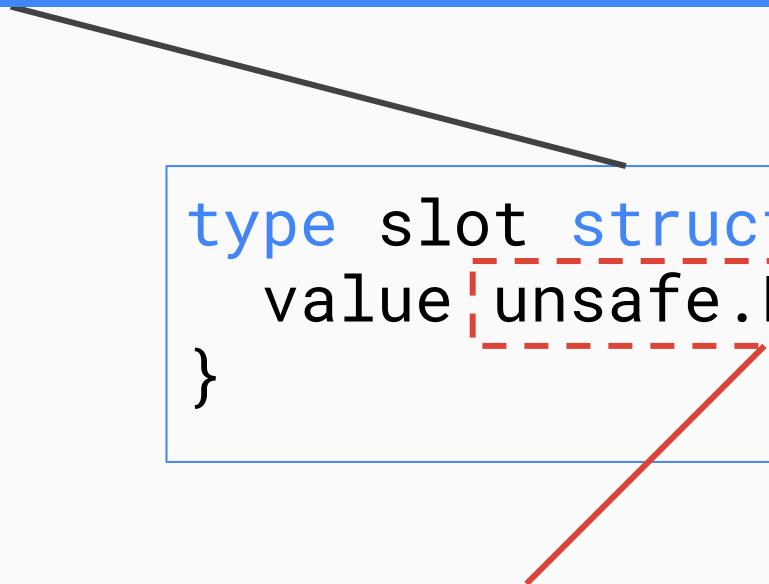
`uintptr` does not retain `pointers` neither

r0

r1

r2

...



```
type slot struct {  
    value unsafe.Pointer  
}
```

Not safe to store **scalars** there!

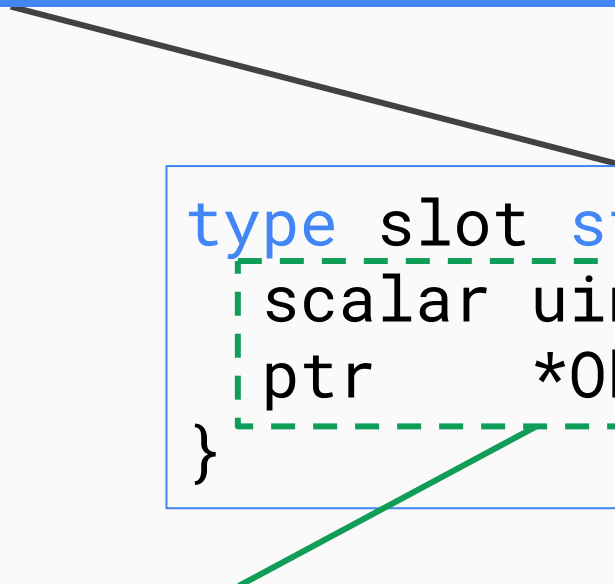
Function frame

r0

r1

r2

...



```
type slot struct {  
    scalar uint64  
    ptr     *Object  
}
```

Paired {scalar, ptr} slots are a safe fix

{uint64, pointer} slots

Function frame

r0

r1

r2

...

scalar

ptr

scalar

ptr

scalar

ptr

Set every slot.ptr to nil

Memory reclaim

Function frame

r0

r1

r2

...

```
type slot struct {  
    value uint64  
}
```

Keeps a
pointer
alive

ρ0

ρ1

ρ2

...

Uint64 + second frame for pointers (alt solution)

Part 4/7

Interop / FFI

- 0 - Backstory
- 1 - go-jdk overview
- 2 - Making the code run fast
- 3 - GC-friendly slots
- > 4 - Interop / FFI
- 5 - Object layout / mem alloc
- 6 - Challenges & limitations
- 7 - Closing words



Calling Java from Go

- Mark machine code buf as executable
- Call as func or do JMP in asm

Simple, boring.

Calling Go from Java

This is more involved (take a breath):

Calling Go from Java

This is more involved (take a breath):

- Obtain Go function address (simple)

Calling Go from Java

This is more involved (take a breath):

- Obtain Go function address (simple)
- Follow the Go calling convention (normal)

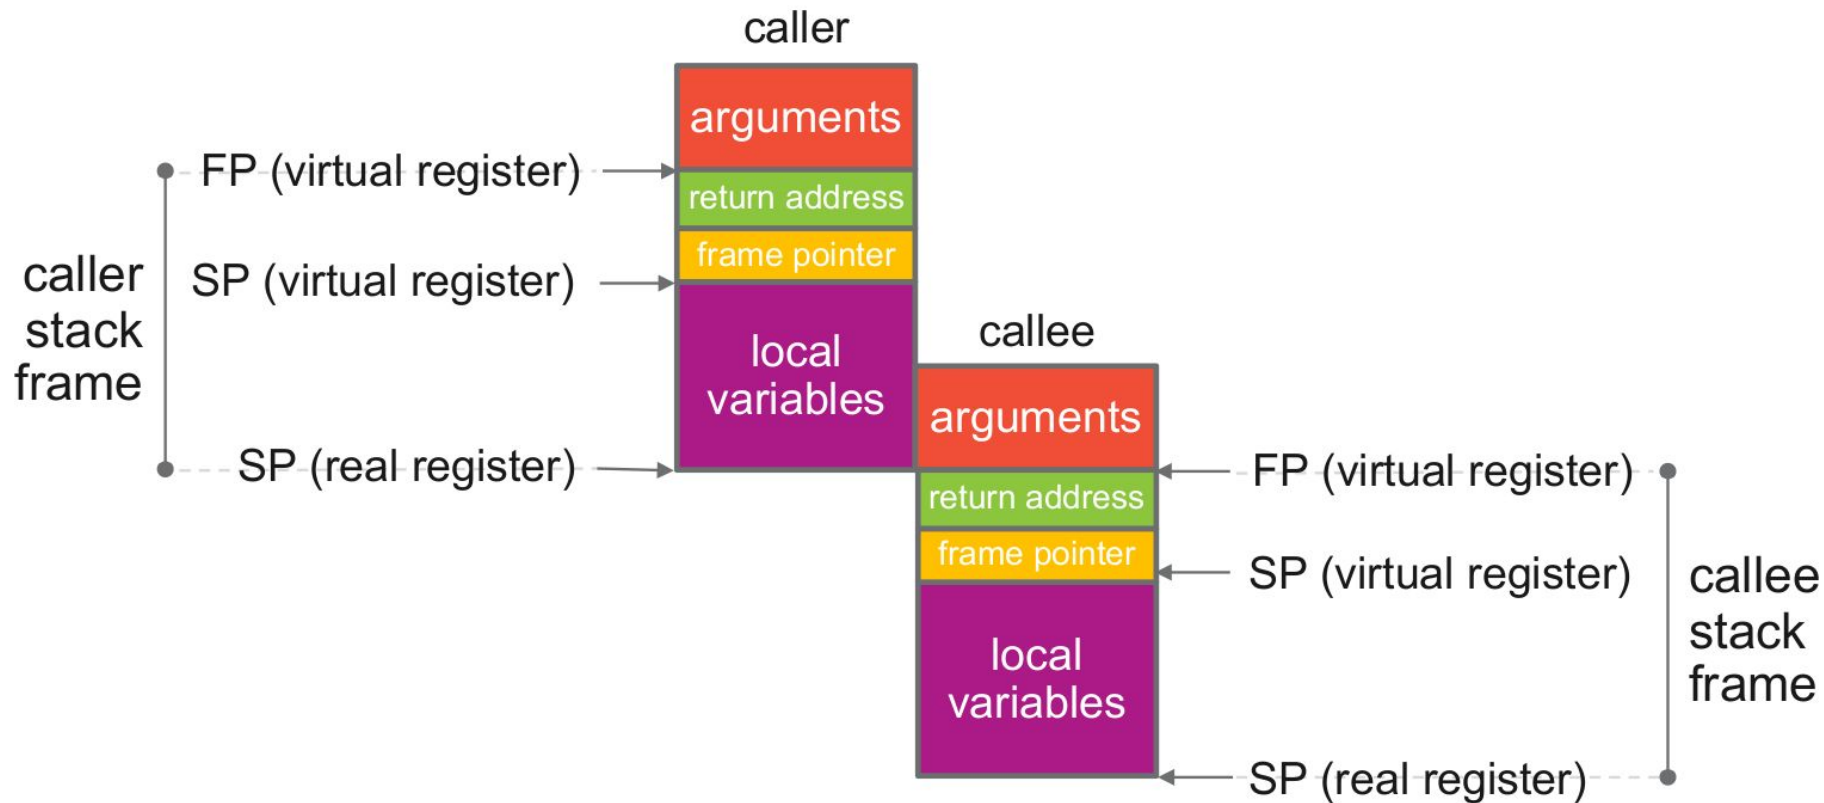
Calling Go from Java

This is more involved (take a breath):

- Obtain Go function address (simple)
- Follow the Go calling convention (normal)
- Deal with fatal error issues (hard)


How to get a Go function code address?

```
func funcAddr(fn interface{}) uintptr {  
    type eface struct {  
        typ    uintptr  
        value *uintptr  
    }  
    e := (*eface)(unsafe.Pointer(&fn))  
    return *e.value  
}
```



Go calling convention ([source](#))


Assembling Java→Go call

1. Push arguments to the stack
2. CALL `$func_addr`  Use funcAddr to get that
3. Move results to local slots

The exact actions depend on the current Go calling convention.

Let's try it!


...




Go runtime is *not* impressed!

- “Unknown caller PC”
- “Unknown return PC”
- “Missing stackmap”

```
=== RUN   TestFoo
runtime: frame <censored> untyped locals 0xc00008ff38+0x8
fatal error: missing stackmap
```





Calling Go directly from the JIT'ed code

You've run into a really hairy area
of asm code.

My first suggestion is not try to
call from assembler into Go.

Ian Lance Taylor

Quote

My first suggestion is not try to
call from assembler into Go.

Ian Lance Taylor

Quote

DON'T UNDERESTIMATE

My first suggestion is not try to
call from assembler into Go.

Ian Lance Taylor

MY POWER

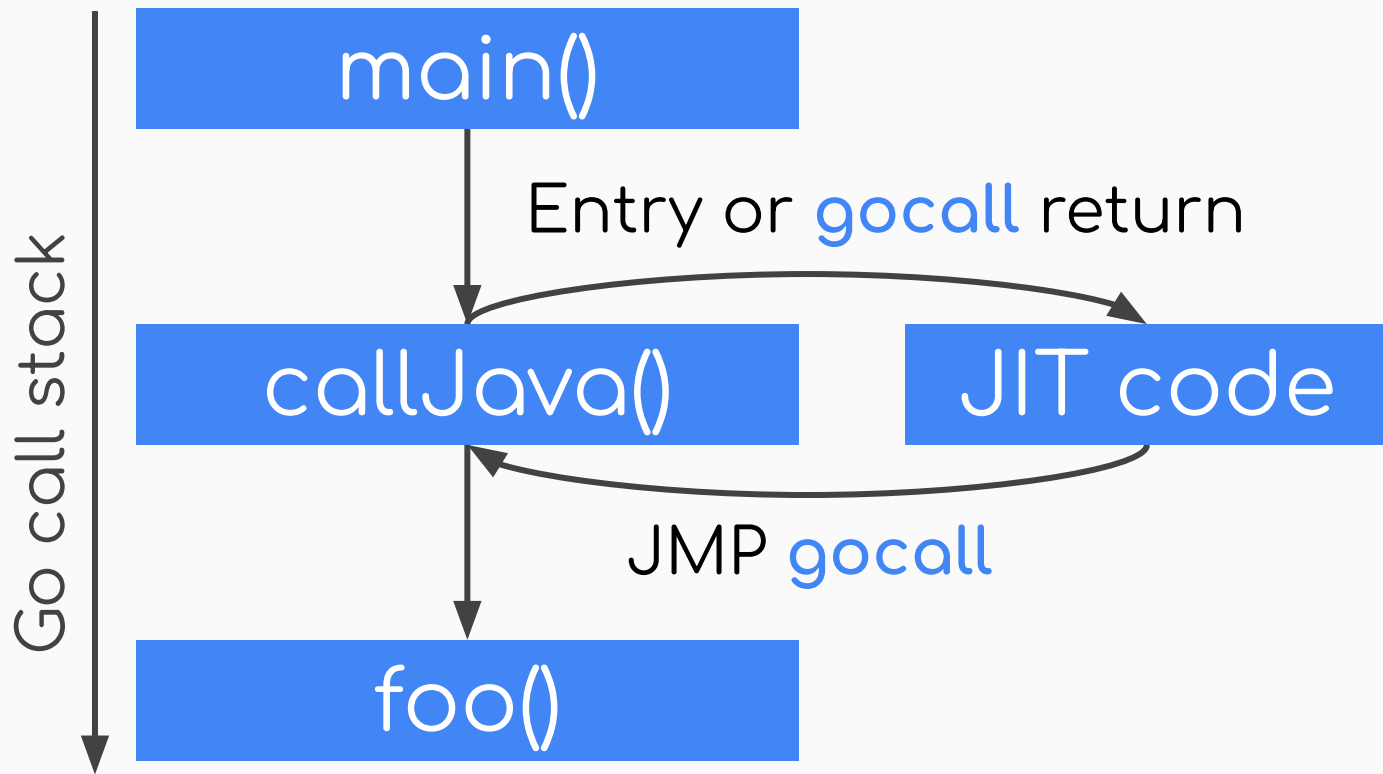
Quote

How to fix these fatals?

Add a Go→Java calls proxy.

Java→Go calls via trampoline.

- Provides a stackmap for Java→Go calls
- Provides a known caller/return PC



Go→Java call proxy (simplified)

```
// callJava(e *Env, code *byte)
TEXT ·callJava(SB), 0, $96-16
    NO_LOCAL_POINTERS
    JMP code+8(FP)
    RET
gocall:
    CALL CX
    JMP -8(BP)
```

Go→Java call proxy (simplified)

```
// callJava(e *Env, code *byte)
```

```
TEXT ·callJava(SB), 0, $96-16
```

```
    NO_LOCAL_POINTERS
```

```
    JMP code+8(FP)
```

```
    RET
```

```
gocall:
```

```
    CALL CX
```

```
    JMP -8(BP)
```



Stackmap fix

NO_LOCAL_POINTERS macro

It's safe for us, as long as:

- We never rely on Go stack values address
- Our heap values are reachable elsewhere

Go→Java call proxy (simplified)

```
// callJava(e *Env, code *byte)
```

```
TEXT ·callJava(SB), 0, $96-16
```

```
    NO_LOCAL_POINTERS
```

```
    JMP code+8(FP)
```

```
    RET
```

```
gocall:
```

```
    CALL CX
```

```
    JMP -8(BP)
```

Caller PC fix

Go→Java call proxy (fixing return PC)

```
// callJava(e *Env, code *byte)
```

```
TEXT ·callJava(SB), 0, $96-16
```

```
    NO_LOCAL_POINTERS
```

```
    MOVQ code+8(FP), CX
```

```
    JCALL(CX)
```

```
    RET
```

```
gocall:
```

```
    CALL CX
```

```
    JMP -8(BP)
```



Return PC fix

Go→Java call proxy (fixing return PC)

```
// callJava(e *Env, code *byte)
```

```
TEXT ·callJava(SB), 0, $96-16
```

```
    NO_LOCAL_POINTERS
```

```
    MOVQ code+8(FP), CX
```

```
    JCALL(CX)
```


```
    RET
```

```
gocall:
```

```
    CALL CX
```

```
    JMP -8(BP)
```

Saves following RET inst
addr and Jumps to CX
(see next slide)



JCALL macro

```
// Encoding `lea rax, [rip+N]` with BYTE
// since Go has no real RIP-relative
// addressing mode.
#define JCALL(fnreg) \
    BYTE $0x48; ... 8d0509000000 \ // Lea
    MOVQ AX, (SI) \ // Store RET addr
    ADDQ $16, SI \ // Move to next slot
    JMP fnreg // Run JIT code
```

Java native methods

```
// In Java file:  
public class Foo {  
    public static native void printInt(int x);  
}
```

Java native methods

```
// In Java file:  
public class Foo {  
    public static native void printInt(int x);  
}  
  
// In Go file:  
func fooPrintInt(x int32) {  
    fmt.Println(x)  
}
```

Java native methods

```
// In Java file:
public class Foo {
    public static native void printInt(int x);
}

// In Go file:
func fooPrintInt(x int32) {
    fmt.Println(x)
}

// Before loading Foo class:
vm.Bind("Foo.printInt", fooPrintInt)
```

Why do we need fast Java→Go?

If calls to Go are fast, we can:

- Implement runtime funcs as Go funcs
- Re-use Go code easily in our Java code

Part 5/7

Object layout and memory allocation

- 0 - Backstory
- 1 - go-jdk overview
- 2 - Making the code run fast
- 3 - GC-friendly slots
- 4 - Interop / FFI
- > 5 - Object layout / mem alloc
- 6 - Challenges & limitations
- 7 - Closing words



Foo class

```
public class Foo {  
    public int x;    // scalar 1  
    public int y;    // scalar 2  
    public Bar bar; // pointer field  
}
```

Foo class

```
public class Foo {  
    public int x;    // scalar 1  
    public int y;    // scalar 2  
    public Bar bar; // pointer field  
}
```

```
type Foo struct {  
    X    int32  
    Y    int32  
    Bar *Bar  
}
```

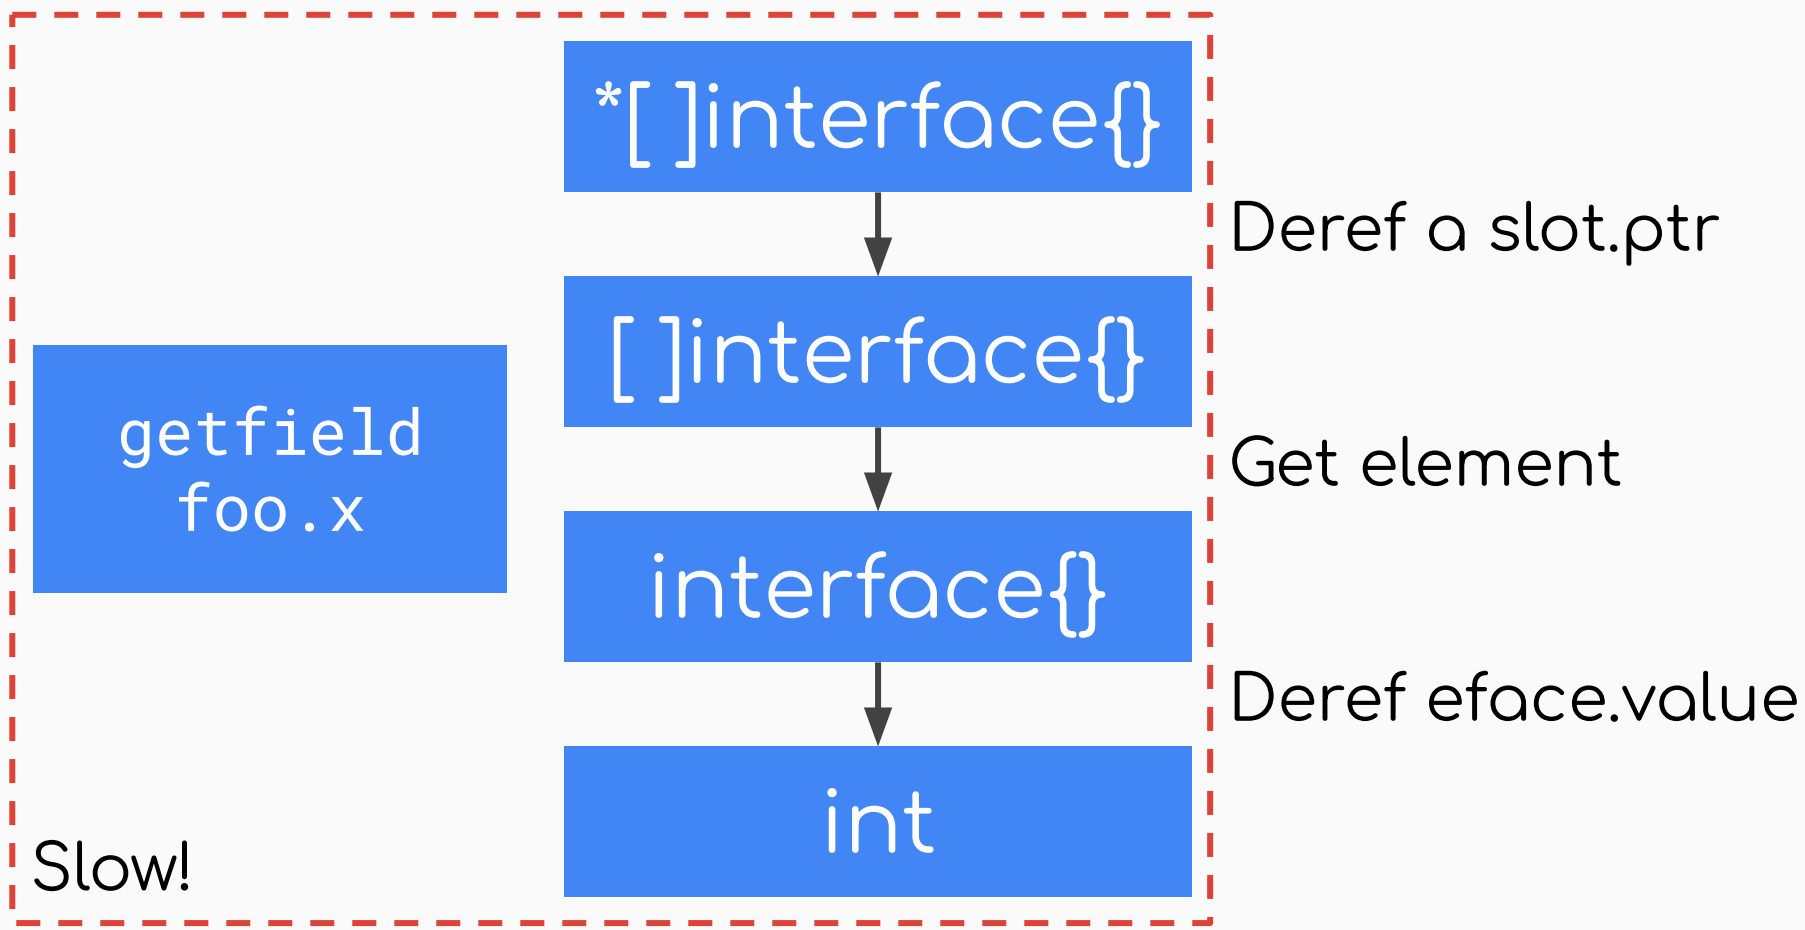
Perfect, but impossible



Foo class values (naive version)

```
// Object is a slice of interface{} fields.  
// Pointer slot gets a slice pointer.
```

```
foo = []interface{}{x, y, bar}  
slot.ptr = &foo
```



Read x:int field from Foo

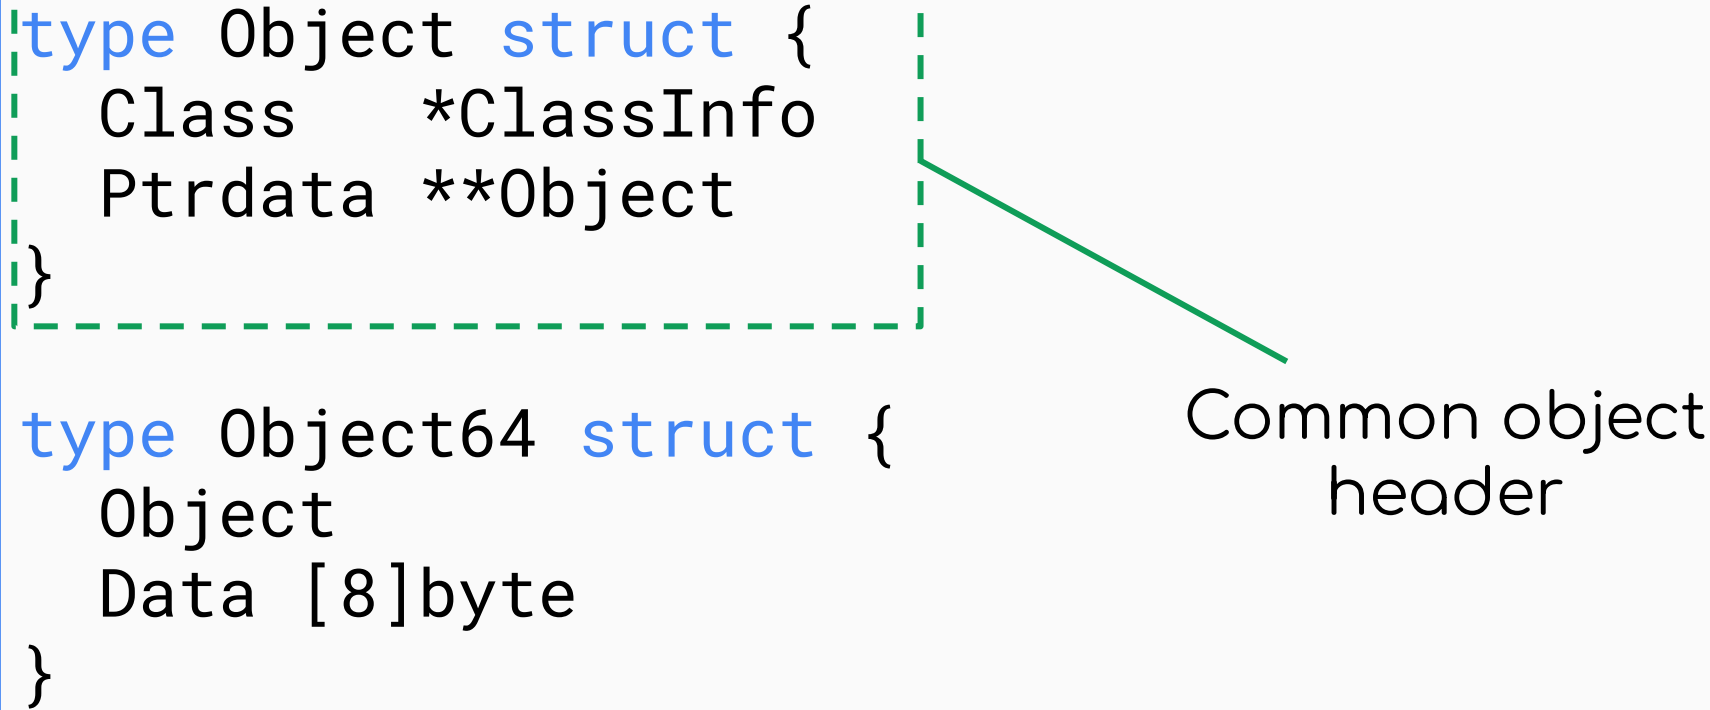
Proposed object layout

```
type Object struct {  
    Class    *ClassInfo  
    Ptrdata **Object  
}
```

```
type Object64 struct {  
    Object  
    Data [8]byte  
}
```

Proposed object layout

```
type Object struct {  
    Class    *ClassInfo  
    Ptrdata  **Object  
}
```



```
type Object64 struct {  
    Object  
    Data [8]byte  
}
```

Common object
header

Proposed object layout

```
type Object struct {  
    Class    *ClassInfo  
    Ptrdata **Object  
}
```

```
type Object64 struct {  
    Object  
    Data [8]byte  
}
```

All object
pointer fields
are stored here



Proposed object layout

```
type Object struct {  
    Class    *ClassInfo  
    Ptrdata  **Object  
}
```

```
type Object64 struct {  
    Object  
    Data [8]byte  
}
```

Object with
8-byte storage
for scalar fields,
Object<64>

Proposed object layout


```
type Object struct {  
    Class    *ClassInfo  
    Ptrdata **Object  
}
```

```
type Object64 struct {  
    Object  
    Data [8]byte  
}
```

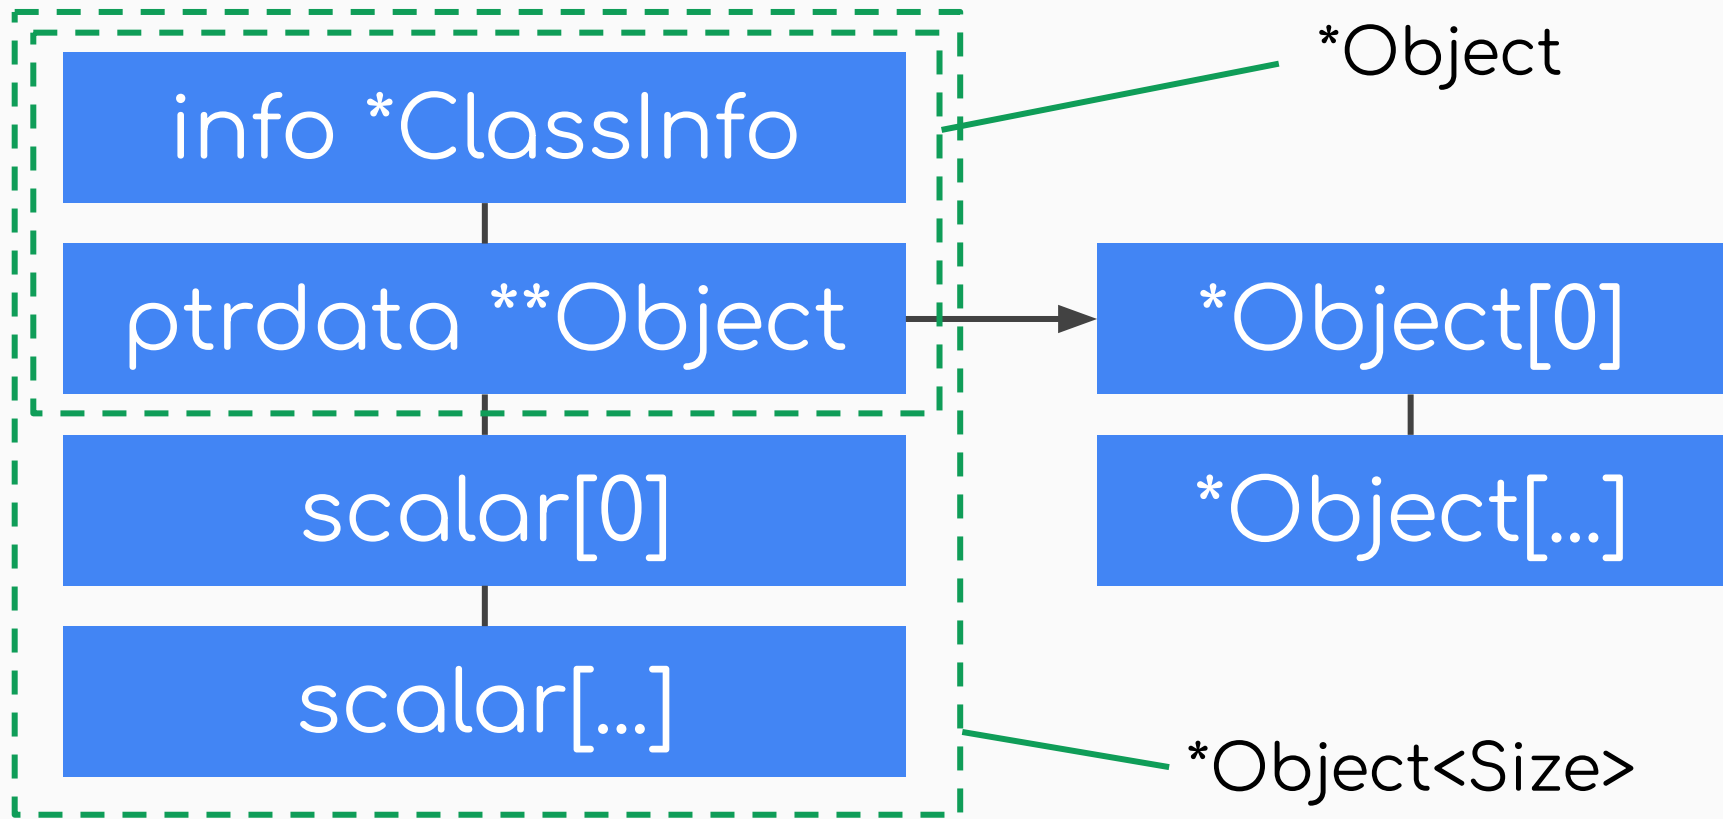
X and Y fields
can be stored
here



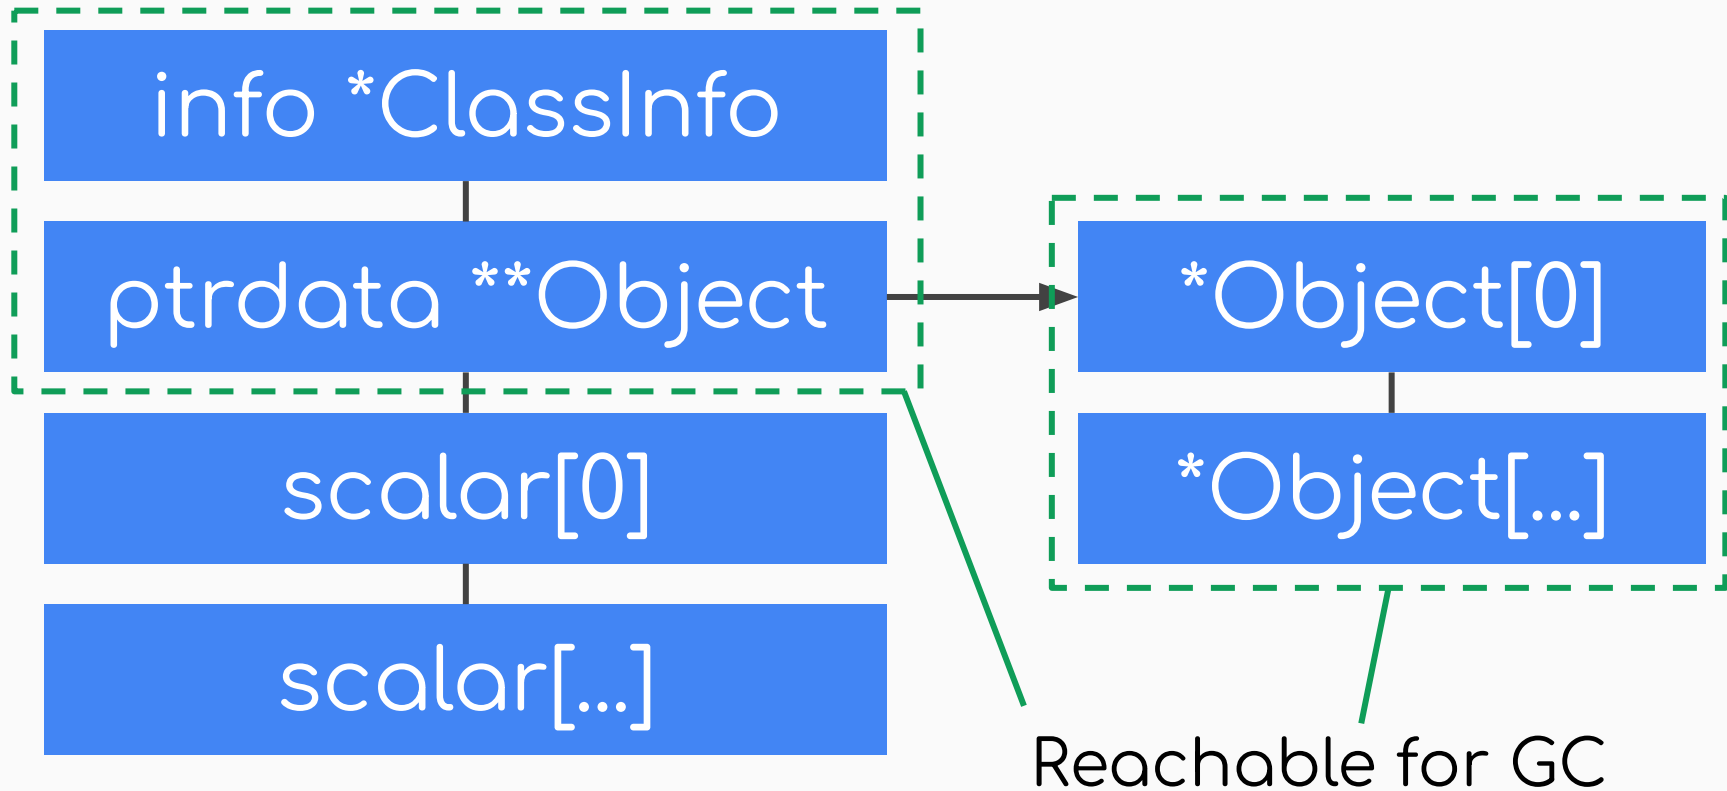
Violates “unsafe”
package rules
(but it’s still OK)



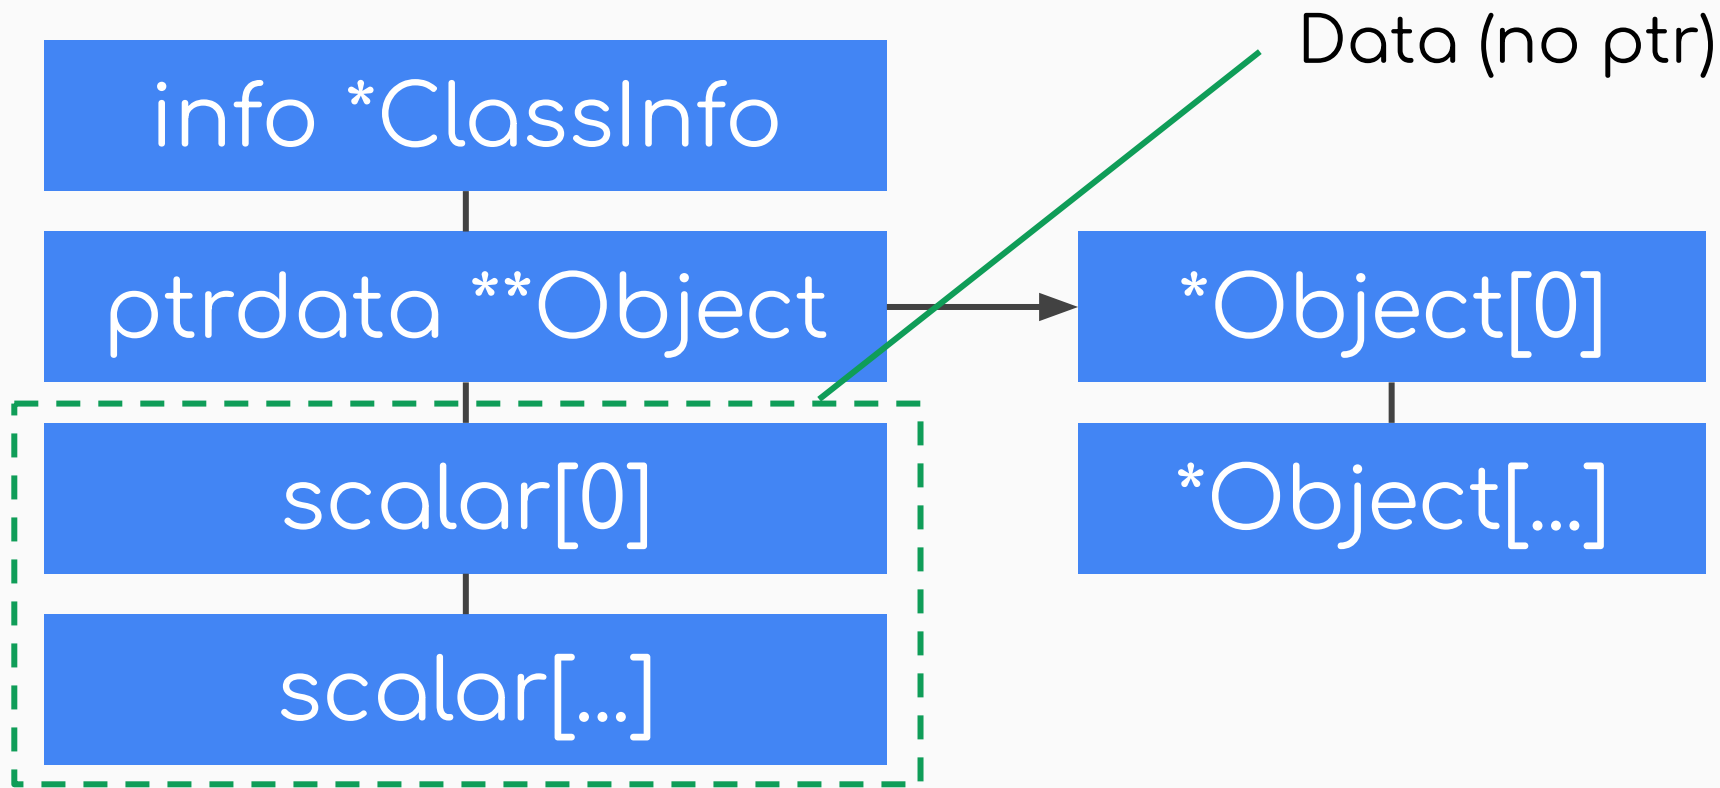
Conversion between `Object` and `Object<Size>`



Abstract Object<Size> layout



Abstract Object<Size> layout



Abstract Object<Size> layout

info *ClassInfo

ptrdata **Object

x:int

y:int

bar:*Object

Foo layout in memory

getfield
foo.x


*Object



int

Deref a slot.ptr
At a proper offset

Read x:int field from Foo



Read bar:Bar field from Foo

Can we use []byte allocations?

No, Go GC will not track any pointers that are stored inside that memory.

So, how to allocate?

- Choose the closest `Object<Size>`
- Allocate `Object<Size>`
- Return as `*Object`

May want to adjust sizes to the Go memory allocator size classes.

For *huge* objects we
can use a less
efficient fallback

Object64

Object [64]byte

Object128

Object [128]byte

Object256

Object [256]byte

...

How many Object<Size> types do we need?

Part 6/7

Challenges and limitations

- 0 - Backstory
- 1 - go-jdk overview
- 2 - Making the code run fast
- 3 - GC-friendly slots
- 4 - Interop / FFI
- 5 - Object layout / mem alloc
- > 6 - Challenges & limitations
- 7 - Closing words



Null pointer check / explicit

```
var p *int // p is nil  
  
println(*p)
```

Null pointer check / explicit

```
var p *int    // p is nil  
nilcheck(p)  // Inserted by a compiler  
println(*p)
```

Null pointer check / explicit


```
var p *int    // p is nil  
nilcheck(p)  // Inserted by a compiler  
println(*p)
```

Simple, but not very efficient

Null pointer check / signals

Hardware exceptions and interrupts
+
OS signals handling

More: <https://stackoverflow.com/a/36955888/4017439>



Remember this picture?

Limitation: bytecode patching

For some reasons, it's quite common in Java world to modify the bytecode that is being loaded...

Limitation: bytecode patching

For some reasons, it's quite common in Java world to modify the bytecode that is being loaded...

Since we convert bytecode into the machine code, we have a problem...

Challenge: method re-load

If method changes and we can't fit its code into the old executable buffer, method address will change...

Challenge: method re-load

If method changes and we can't fit its code into the old executable buffer, method address will change...

This requires re-linking all method callers.
If calls were inlined it's even harder.

Part 7/7

Closing words

- 0 - Backstory
- 1 - go-jdk overview
- 2 - Making the code run fast
- 3 - GC-friendly slots
- 4 - Interop / FFI
- 5 - Object layout / mem alloc
- 6 - Challenges & limitations
- > 7 - Closing words



Testing

```
import testutil.T;
```

```
class Test {  
    public static void run(int x) {  
        T.println(x + 5);  
    }  
}
```



System.out.println in OpenJDK,
fmt.Println in go-jdk

N-body benchmark results

OpenJDK	3.9s
go-jdk	4.8s
OpenJDK (no JIT)	~11s
go-jdk (no JIT)	~22s

N-body benchmark results

OpenJDK	3.9s
go-jdk	4.8s
OpenJDK (no JIT)	~11s
go-jdk (no JIT)	~22s

N-body benchmark results

OpenJDK	3.9s
go-jdk	4.8s
OpenJDK (no JIT)	~11s
go-jdk (no JIT)	~22s

Resources

- [go-jdk repository](#)
- [VM Showdown: Stack Versus Registers](#)
- [Calling Go funcs from asm](#) (ru)
- [Go calling convention](#)
- [JNI bindings for Go](#)

Efficient VM with JIT in Go

quasilyte @ GoWayFest 4.0 (2020)

