

# Deterministic performance tests

quasilyte @ GDG Nizhny Novgorod  
Go meetup, 2021

# Measurable performance attributes

# Measurable performance attributes

- Execution time

# Measurable performance attributes

- Execution time
- Memory consumption

# Measurable performance attributes

- Execution time
- Memory consumption
- Heap allocations / GC pressure

# Non-deterministic ways

# Non-deterministic ways

- Benchmark tests

# Non-deterministic ways

- Benchmark tests
- Profiling (of any kind)



# Non-deterministic ways

- Benchmark tests
- Profiling (of any kind)
- Metrics, logs, traces (allocs, rps, etc.)

# A conventional work-around

1. Run benchmarks on a dedicated machine
2. Save benchmark results from various revisions
3. Compare new rev result with older results

The “test” is failed if the new code gave worse results than we recorded before.

HEAD~2

Bench  
results  
rev~f990ca


HEAD~1

Bench  
results  
rev~129ab4


HEAD

Bench  
results  
rev~cfd105

Comparing benchmark results



Comparing benchmark results



Comparing benchmark results

What is a “deterministic perf test”?

# Defining a deterministic performance test

# Defining a deterministic performance test

- Results are stable



# Defining a deterministic performance test

- Results are stable
- Results are reproducible on different machines

# Defining a deterministic performance test

- Results are stable
- Results are reproducible on different machines

But, more importantly...

# Defining a deterministic performance test

- Results are stable
- Results are reproducible on different machines

But, more importantly...

- They **fail** if their invariant is broken

# What we can measure reliably

- Number of the computational steps
- Heap allocations
- Whether Go compiler applied some optimization

# What we can measure reliably

- Number of the computational steps
- Heap allocations
- Whether Go compiler applied some optimization

Safe zone

# What we can measure reliably

- Number of the computational steps
- Heap allocations
- Whether Go compiler applied some optimization

Gray zone!

Test may be more fragile than you want it to be

# What we can measure reliably

- Number of the computational steps
- Heap allocations
- Whether Go compiler applied some optimization

Danger zone!

Hard Go toolchain version dependency

# The universal solution



# The universal solution

Doesn't exist

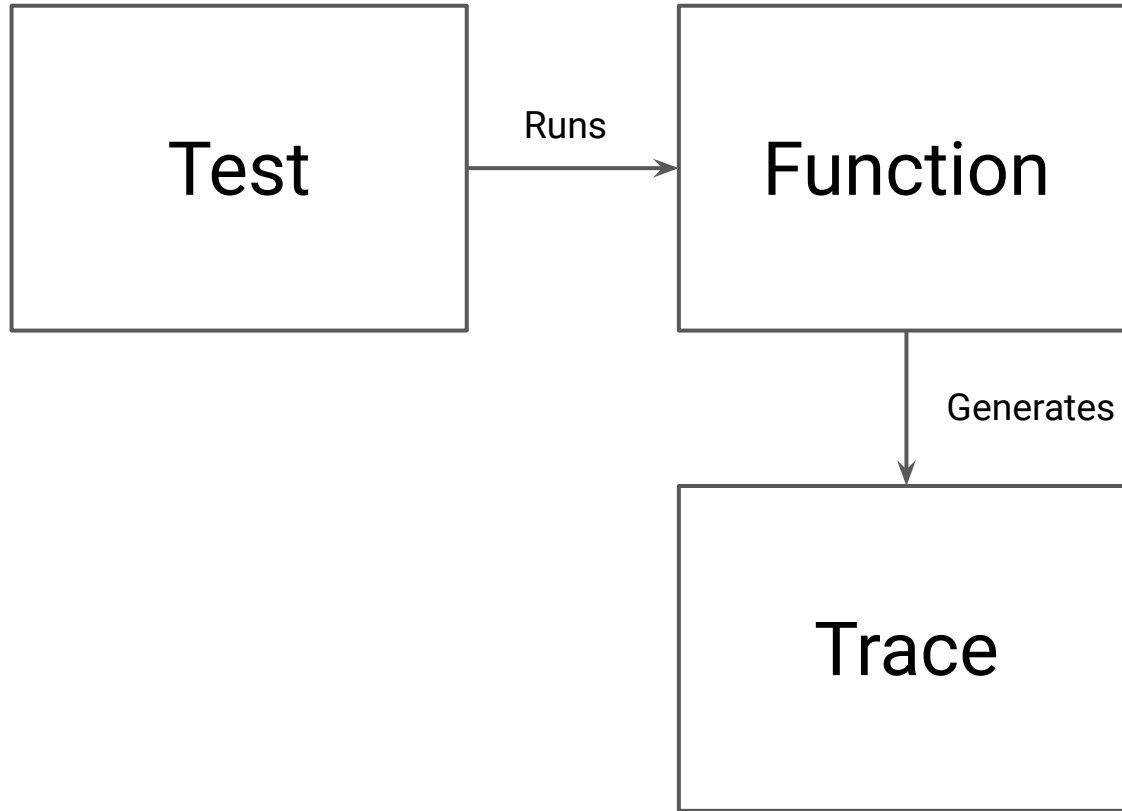
The universal solution

Doesn't exist


But I know some magic tricks

# Magic trick #1

Execution trace tests



Execution trace tests



Execution trace tests

# Execution trace tests

```
// When outer matching fails,  
// no other matching should be performed.  
{  
  pattern: `if (0) $_`,  
  input:   `if (1) { if (2) { if (3) {} } }`,  
  trace: []string{  
    "eqNode x=*ir.IfStmt y=*ir.IfStmt",  
    " • eqNode x=*ir.Lnumber y=*ir.Lnumber",  
  },  
}
```

# Execution trace tests

```
// When outer matching fails,  
// no other matching should be performed.  
{  
  pattern: `if (0) $_`,  
  input:   `if (1) { if (2) { if (3) {} } }`,  
  trace: []string{  
    "eqNode x=*ir.IfStmt y=*ir.IfStmt",  
    " • eqNode x=*ir.Lnumber y=*ir.Lnumber",  
  },  
}
```

# Execution trace tests

```
// When outer matching fails,  
// no other matching should be performed.  
{  
  pattern: `if (0) $_`,  
  input:   `if (1) { if (2) { if (3) {} } }`,  
  trace: []string{  
    "eqNode x=*ir.IfStmt y=*ir.IfStmt",  
    " • eqNode x=*ir.Lnumber y=*ir.Lnumber",  
  },  
}
```



# Execution trace tests

```
// When $_ is used for statement body it should quickly
// match the body without checking the actual contents.
{
  pattern: `if ($cond) $_`,
  input:   `if ($x && g()) { echo 1, 2, 3; }`,
  trace: []string{
    "eqNode x=*ir.IfStmt y=*ir.IfStmt",
    " • eqNode x=*ir.SimpleVar y=*ir.BooleanAndExpr",
    " • eqNode x=*ir.ExpressionStmt y=*ir.StmtList",
  },
},
```

# Execution trace tests

```
// When $_ is used for statement body it should quickly
// match the body without checking the actual contents.
{
  pattern: `if ($cond) $_`,
  input:   `if ($x && g()) { echo 1, 2, 3; }`,
  trace: []string{
    "eqNode x=*ir.IfStmt y=*ir.IfStmt",
    " • eqNode x=*ir.SimpleVar y=*ir.BooleanAndExpr",
    " • eqNode x=*ir.ExpressionStmt y=*ir.StmtList",
  },
},
```

# Execution trace tests

```
// When $_ is used for statement body it should quickly
// match the body without checking the actual contents.
{
  pattern: `if ($cond) $_`,
  input:   `if ($x && g()) { echo 1, 2, 3; }`,
  trace: []string{
    "eqNode x=*ir.IfStmt y=*ir.IfStmt",
    " • eqNode x=*ir.SimpleVar y=*ir.BooleanAndExpr",
    " • eqNode x=*ir.ExpressionStmt y=*ir.StmtList",
  },
},
```

# Execution trace tests

```
// When $_ is used for statement body it should quickly
// match the body without checking the actual contents.
{
  pattern: `if ($cond) $_`,
  input:   `if ($x && g()) { echo 1, 2, 3; }`,
  trace: []string{
    "eqNode x=*ir.IfStmt y=*ir.IfStmt",
    " • eqNode x=*ir.SimpleVar y=*ir.BooleanAndExpr",
    " • eqNode x=*ir.ExpressionStmt y=*ir.StmtList",
  },
},
```

# Execution trace tests

```
// Match $_ once, then accept match as ${"*"} is the last node.
{
  pattern: `f($_, ${"*"})`,
  input:   `f(1, 2, 3, 4, g())`,
  trace: []string{
    "eqNode x=*ir.FunctionCallExpr y=*ir.FunctionCallExpr",
    " • eqNode x=*ir.Argument y=*ir.Argument",
    " • • eqNode x=*ir.SimpleVar y=*ir.Lnumber",
  },
},
```

# Execution trace tests

```
// Match $_ once, then accept match as ${"*"} is the last node.
{
  pattern: `f($_, ${"*"})`,
  input:   `f(1, 2, 3, 4, g())`,
  trace: []string{
    "eqNode x=*ir.FunctionCallExpr y=*ir.FunctionCallExpr",
    " • eqNode x=*ir.Argument y=*ir.Argument",
    " • • eqNode x=*ir.SimpleVar y=*ir.Lnumber",
  },
},
```

# Execution trace tests

```
// Match $_ once, then accept match as ${"*"} is the last node.
{
  pattern: `f($_, ${"*"})`,
  input:   `f(1, 2, 3, 4, g())`,
  trace: []string{
    "eqNode x=*ir.FunctionCallExpr y=*ir.FunctionCallExpr",
    " • eqNode x=*ir.Argument y=*ir.Argument",
    " • • eqNode x=*ir.SimpleVar y=*ir.Lnumber",
  },
},
```

# Execution trace tests

```
// Match $_ once, then accept match as ${"*"} is the last node.
{
  pattern: `f($_, ${"*"})`,
  input:   `f(1, 2, 3, 4, g())`,
  trace: []string{
    "eqNode x=*ir.FunctionCallExpr y=*ir.FunctionCallExpr",
    " • eqNode x=*ir.Argument y=*ir.Argument",
    " • • eqNode x=*ir.SimpleVar y=*ir.Lnumber",
  },
},
```



# Making the traces “free”

1. Collect traces under if with const condition
2. Set const value via the compilation tag

Trace collector struct can be empty for release builds and a slice of records for test runs.

tracing\_enabled.go

```
// +build tracing
```

```
package phpgrep
```

```
const tracingEnabled = true
```

# tracing\_enabled.go

```
// +build tracing
```



```
package phpgrep
```

Or with a newer “go:build”  
syntax

```
const tracingEnabled = true
```

tracing\_disabled.go

```
// +build !tracing
```

```
package phpgrep
```

```
const tracingEnabled = false
```

# Matcher struct

```
type matcher struct {  
    // . . . other fields  
  
    // Used only when -tracing build tag is specified.  
    tracingBuf    *bytes.Buffer  
    tracingDepth  int  
}
```

# Writing the trace

```
if tracingEnabled {  
    pad := strings.Repeat(" • ", m.tracingDepth)  
    fmt.Fprintf(m.tracingBuf, "%seqNode x=%T y=%T\n", pad, x, y)  
    m.tracingDepth++  
    defer func() {  
        m.tracingDepth--  
    }()  
}
```

# Writing the trace

```
if tracingEnabled {  
    pad := strings.Repeat(" • ", m.tracingDepth)  
    fmt.Fprintf(m.tracingBuf, "%seqNode x=%T y=%T\n", pad, x, y)  
    m.tracingDepth++  
    defer func() {  
        m.tracingDepth--  
    }()  
}
```

Const condition:

If tracingEnabled=false, this branch is removed

Where applicable



# Where applicable

- Any kinds of state machines

# Where applicable

- Any kinds of state machines
- Data transformers

# Magic trick #2

Codegen tests

# Codegen kinds (simplified!)

## Codegen kinds (simplified!)

- Ad-hoc (go:generate, codegen scripts)

# Codegen kinds (simplified!)

- Ad-hoc (go:generate, codegen scripts)
- Compilers-like software

# Why testing the codegen?

# Why testing the codegen?

- Avoid unexpected perf regressions



# Why testing the codegen?

- Avoid unexpected perf regressions
- Keep codegen optimization promises

# Why testing the codegen?

- Avoid unexpected perf regressions
- Keep codegen optimization promises
- You'll probably need these tests anyway

# Codegen test: compilers

```
{  
    `if b { return 1 }; return 0`: {  
        ` PushParam 2 # b`,  
        ` JumpFalse 6 # L0`,  
        ` PushIntConst 0 # value=1`,  
        ` ReturnIntTop`,  
        `L0:`,  
        ` PushIntConst 1 # value=0`,  
        ` ReturnIntTop`,  
    },  
    // ...  
}
```

# Codegen test: compilers

```
{  
  `if b { return 1 }; return 0`:{  
    `  PushParam 2 # b`,  
    `  JumpFalse 6 # L0`,  
    `  PushIntConst 0 # value=1`,  
    `  ReturnIntTop`,  
    `L0:`,  
    `  PushIntConst 1 # value=0`,  
    `  ReturnIntTop`,  
  },  
  // ...  
}
```

Input (source)

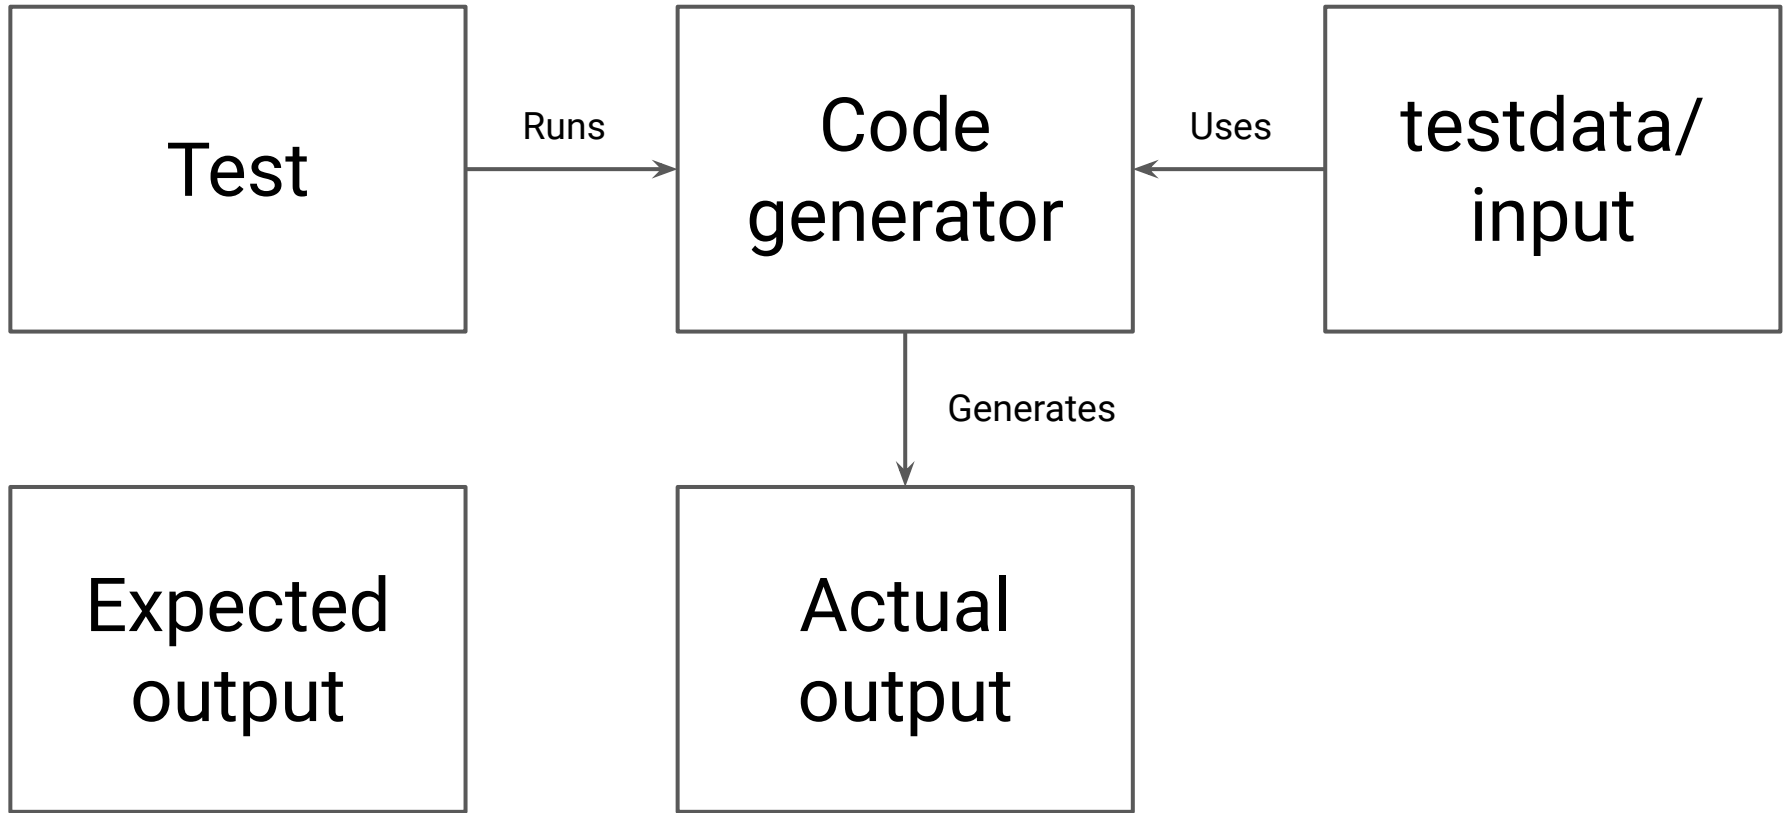
# Codegen test: compilers

```
{  
    `if b { return 1 }; return 0`: {  
        `    PushParam 2 # b`,  
        `    JumpFalse 6 # L0`,  
        `    PushIntConst 0 # value=1`,  
        `    ReturnIntTop`,  
        `L0:`,  
        `    PushIntConst 1 # value=0`,  
        `    ReturnIntTop`,  
    },  
    // ...  
}
```

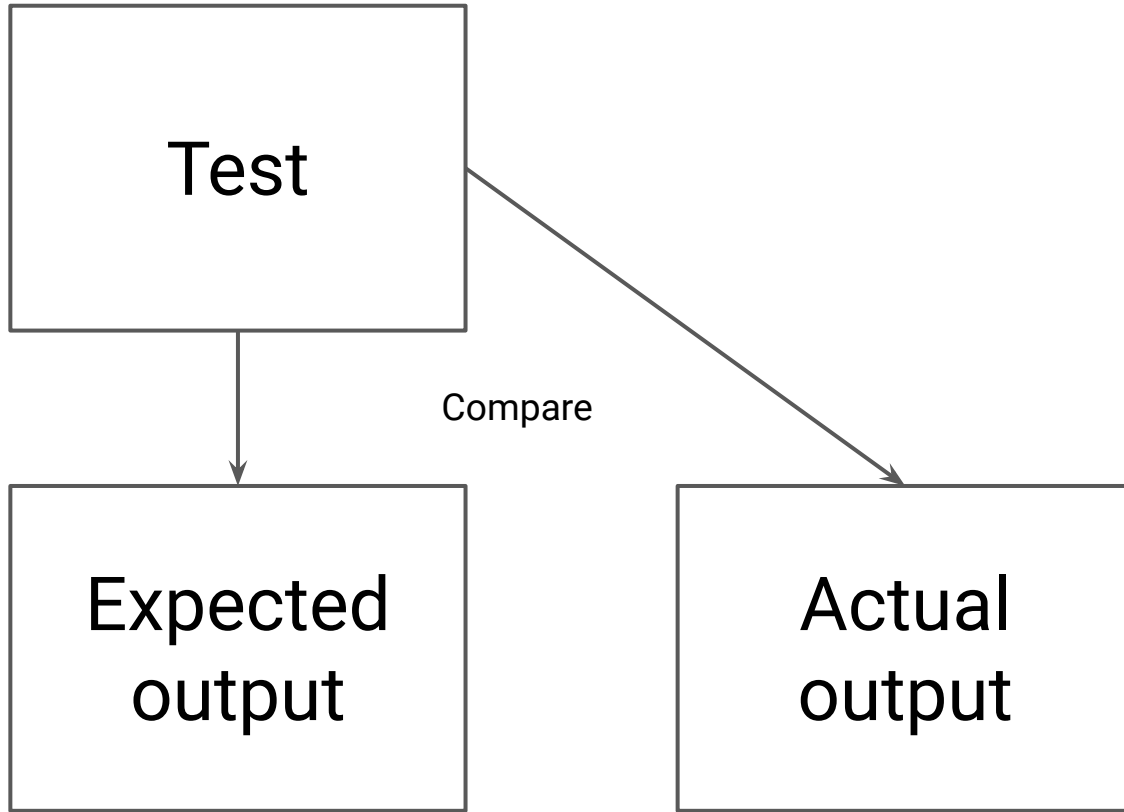
Output (generated code)

# Codeget test: ad-hoc

1. Put input files to “testdata/”
2. Put “expected result” files somewhere close
3. Run codegen over a testdata folder
4. Compare actual and expected results



Ad-hoc codegen test



Ad-hoc codegen test



Where applicable

# Where applicable

- ORM-like code generators

# Where applicable

- ORM-like code generators
- Any auto-generated perf sensitive code

# Where applicable

- ORM-like code generators
- Any auto-generated perf sensitive code
- Compilers, interpreters, optimizers

# Magic trick #3

Zero alloc tests

# Zero alloc tests

1. Collect mem stats (“before”)
2. Run the code to be tested
3. Collect mem stats again (“after”)
4. Compare “before” and “after” allocs

# Zero alloc test

```
var before, after runtime.MemStats
runtime.GC()

runtime.ReadMemStats(&before)
funcToTest() // <- HERE
runtime.ReadMemStats(&after)

allocated := after.Alloc - before.Alloc
if allocated != 0 {
    t.Error(...)
}
```

# Problem

This method can be flaky\*

(\*) fails from time to time



# Solution

Run the test several times

If it `f()` allocates, it always allocates

# Zero alloc test (improved)

```
const numTests = 5 // Five is good enough, 10 is even better
failures := 0

for i := 0; i < numTests; i++ {
    // ... test body.

    if allocated != 0 {
        failures++
    }
}

if failures == numTests {
    t.Error(...)
}
```

Where applicable

# Where applicable

- Low-level functions

# Where applicable

- Low-level functions
- High-performance primitives

# Where applicable

- Low-level functions
- High-performance primitives
- Latency-sensitive code paths

# Magic trick #4

Sizeof tests

# ruleguard AST matching engine

```
type instruction struct {  
    op          operation  
    value       uint8  
    valueIndex  uint8  
}
```

```
type program struct {  
    insts    []instruction  
    strings  []string  
    ifaces   []interface{}  
}
```



# ruleguard AST matching engine

```
type instruction struct {  
    op          operation  
    value       uint8  
    valueIndex  uint8  
}
```

```
type program struct {  
    insts  []instruction  
    strings []string  
    ifaces []interface{}  
}
```

Should be small to stay  
cache-friendly

# Sizeof test

```
wantSize := 3
haveSize := int(unsafe.Sizeof(instruction{}))
if wantSize != haveSize {
    t.Errorf("sizeof(instruction): have %d, want %d",
        haveSize, wantSize)
}
```

# Problem

Sizeof is platform-dependent

# Solution 1

Check for the platform and/or int size

# Sizeof test: adjusting expected sizes

```
is64bit := bits.UintSize == 64
wantSize := 12
if is64bit {
    wantSize = 24
}
```

# Solution 2

Skip platforms that you don't care about

## Sizeof test: skipping platforms

```
if bits.UintSize != 64 {  
    t.Skip("not 64-bit platform")  
}
```

Where applicable



# Where applicable

- Core data points

# Where applicable

- Core data points
- Layout-sensitive data

# Magic trick #5

Go compiler optimizations test

# Disclaimer

You probably don't need to go this far

# What we may want to test

- Is function `f()` inlineable?
- Does function `f()` contain a bound check?
- Do `f()` params/results escape to heap?

# What we may want to test

- Is function `f()` inlineable?
- Does function `f()` contain a bound check?
- Do `f()` params/results escape to heap?

Go compiler can provide answers to all these questions if we'll just ask.

# Algorithm

- Run Go compiler with “-m -m” flags
- Collect the output
- Parse output for the relevant info

# Algorithm

- Run Go compiler with “-m -m” flags
- Collect the output
- Parse output for the relevant info

Note that the output format is not stable between the Go versions.



## inltest library

A legacy library that probably doesn't work anymore, but it can be used as a reference:

<https://github.com/quasilyte/inltest>

Where applicable

# Where applicable

- Anywhere if you're a crazy person

Don't just run benchmarks

# Don't just run benchmarks

Write performance tests

# Deterministic performance tests

quasilyte @ GDG Nizhny Novgorod  
Go meetup, 2021